

► Virtual Memory

Dr. Manmath N. Sahoo
Dept. of CSE, NIT Rourkela

Virtual Memory

- ▶ **Virtual Memory** is a technique that allows the execution of processes that are not completely in main memory.
- ▶ **Motivation**
 - ▶ Conditional execution:
 - either 2 pages of **if** block will be loaded or 2 pages of **else** block.
 - ▶ Error handling codes in programs are seldom executed.
 - ▶ Even though all the pages are needed, they are not needed at the same time.


```
If(condition){  
    [ ]  
    [ ]  
}  
Else{  
    [ ]  
    [ ]  
}
```

Virtual Memory


► Advantages:

- A process would no longer be constrained by the amount of available physical memory.
- Increases degree of multi-programming.
- Better CPU utilization and throughput.
- Less I/O overhead

Virtual Memory That is Larger Than Physical Memory



Virtual-address Space




Demand Paging

- ▶ Bring a page into memory only when it is needed.
- ▶ Page is needed \Rightarrow reference to it
 - ▶ invalid reference (not-in-memory) \Rightarrow bring to memory

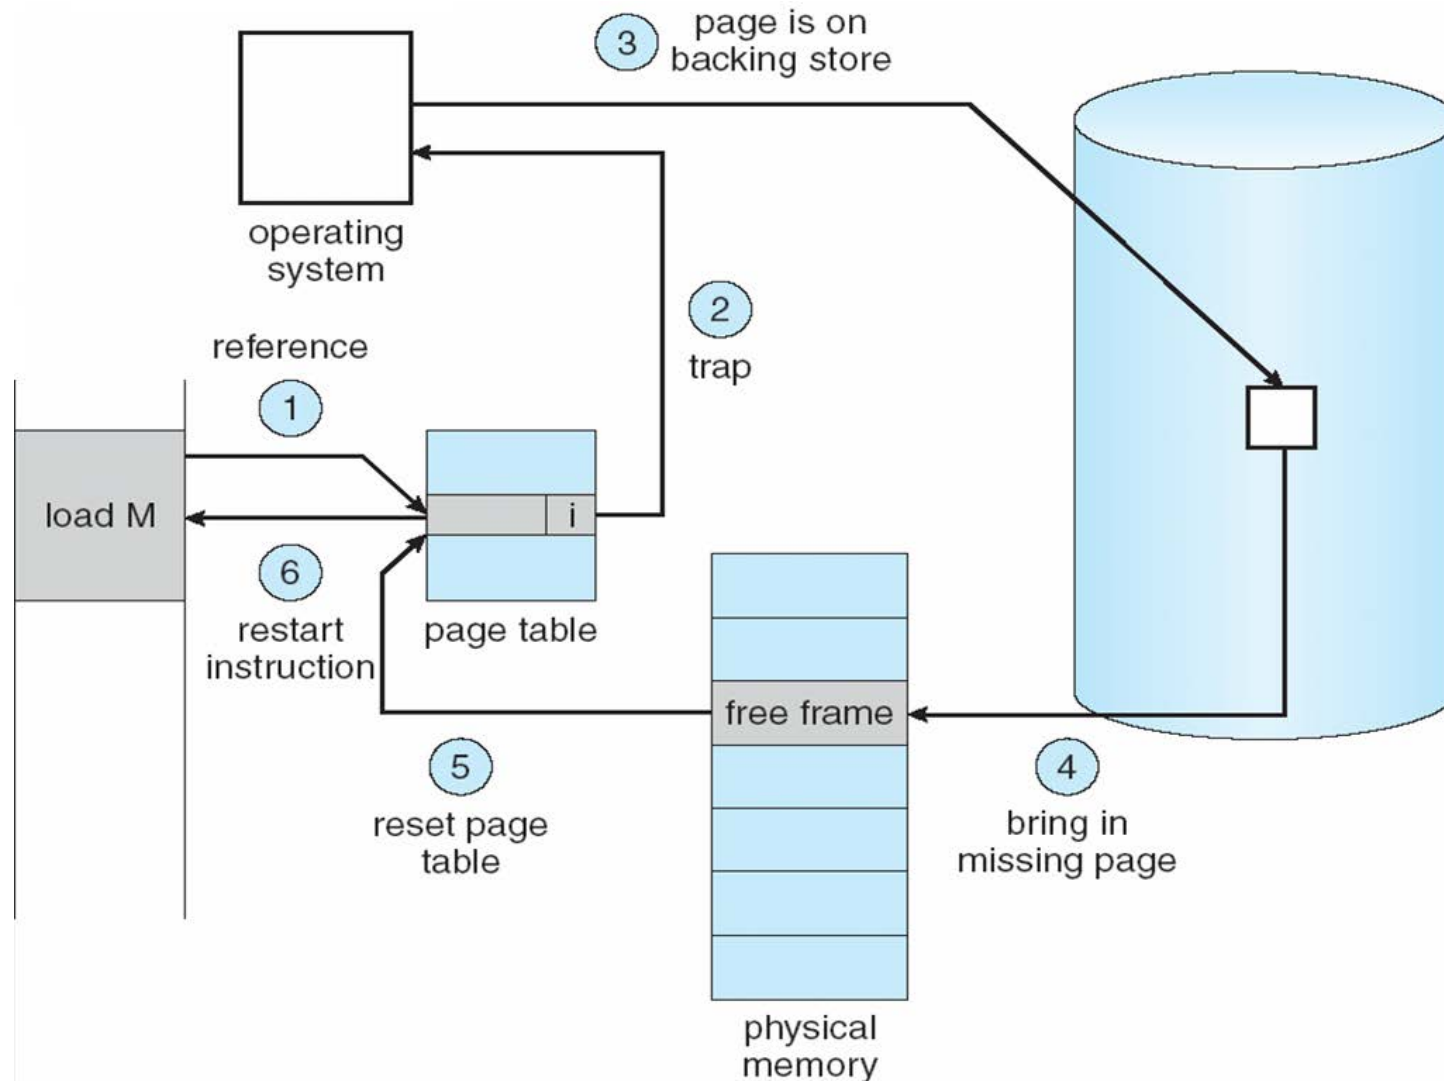
Valid-Invalid Bit

- ▶ With each page table entry a valid–invalid bit is associated (**1** \Rightarrow in-memory, **0** \Rightarrow not-in-memory)
- ▶ Initially valid–invalid bit is set to **0** on all entries
- ▶ During address translation, if valid–invalid bit in page table entry is **0** \Rightarrow page fault

Page Table When Some Pages Are Not in Main Memory



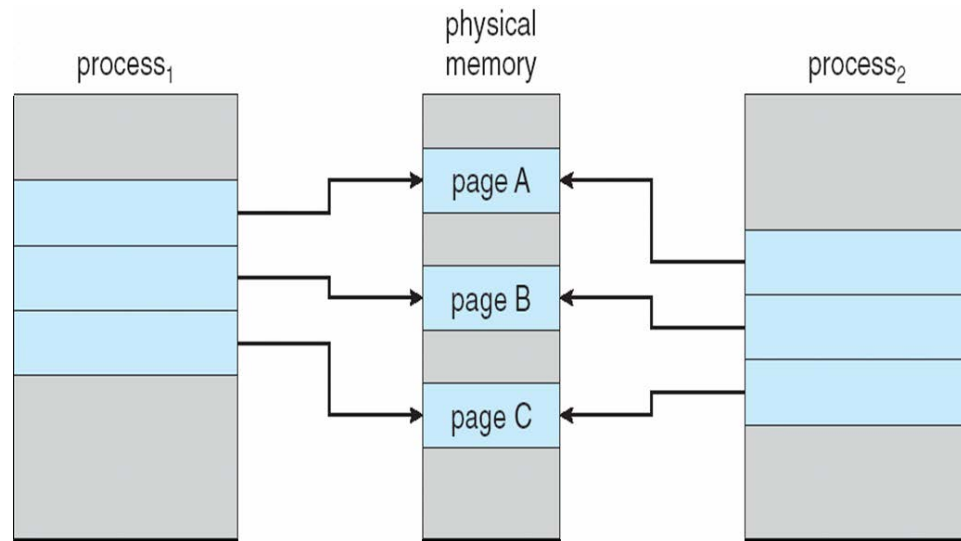
Steps in Handling a Page Fault



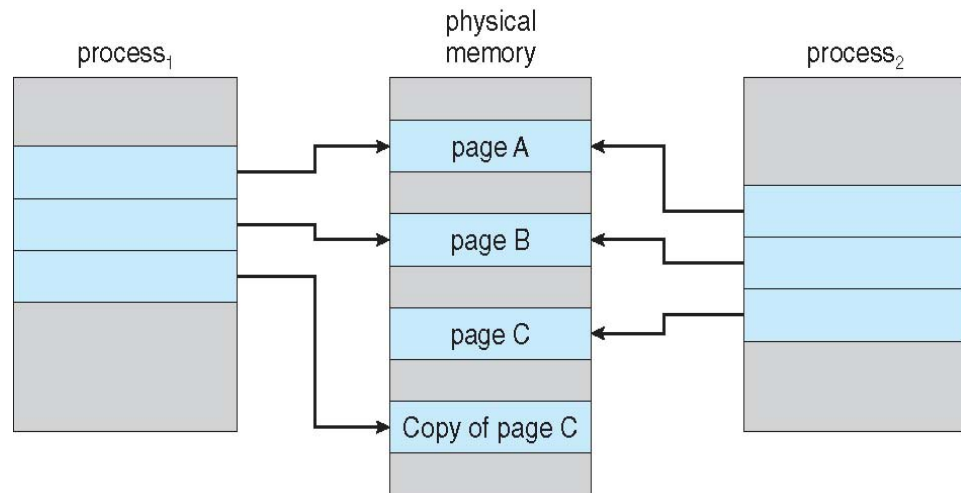
Copy-on-Write

- ▶ Copy-on-Write allows both parent and child processes to initially *share* the same pages in memory
- ▶ If either process modifies a shared page, only then is the page copied
- ▶ It allows more efficient process creation as only modified pages are copied

Copy-on-Write



Before Process 1 Modifies Page C




After Process 1 Modifies Page C

What happens if there is no free frame?

- ▶ Page replacement – find some page in memory, but not really in use, swap it out
 - ▶ algorithm
 - ▶ performance – want an algorithm which will result in minimum number of page faults
 - ▶ Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk

Page Replacement



First-In-First-Out (FIFO) Algorithm

Replace page that has come first.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																
	0	0	0																
		1	1																

2	2	4	4	4	0														
3	3	3	2	2	2														
1	0	0	0	3	3														

0	0																		
1	1																		
3	2																		

7	7	7																	
1	0	0																	
2	2	1																	

page frames

Total page faults = 15

FIFO Algorithm: Belady's Anomaly


- ▶ Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- ▶ 3 frames – 9 page faults
- ▶ 4 frames – 10 page faults

Optimal Algorithm

- Replace page that will not be used for longest period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames


- How do you know this?
- Used for measuring how well your algorithm performs

Least Recently Used (LRU) Algorithm

- Replace the page that has not been used for longest period of time.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

LRU Implementation

▶ Counter implementation

- ▶ Every page entry has a counter; every time page is referenced, copy the clock into the counter
- ▶ Logical counter ticks with every page reference
- ▶ Replace a valid page with smallest counter


LRU Implementation

► Modified Stack implementation


- Push at top only, Pop from anywhere
- Page referenced \Rightarrow on the memory \Rightarrow move it to the top
- LRU page will be at the bottom always

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a



stack
after
b



Additional Reference bits Algorithm

- ▶ Keep a reference/use bit and 7-bit reference history for each entry in PMT.
- ▶ For every page reference, update the reference bit
- ▶ At regular intervals (say 50ns)
 - ▶ Perform right shift to the history bits, discarding the LSB
 - ▶ Push reference bit value to MSB of history bits
 - ▶ Reset reference bit to 0
- ▶ Thus 1-byte represents the status of the page usage in 8 intervals.
 - ▶ 00000000 – The page has not been used for 8 intervals
 - ▶ 11111111 – The page is used at least once in each interval
 - ▶ 11000100 is more recently used than 01110111
 - ▶ So, lowest number is the LRU page

Additional Reference bits Algorithm

3, 2, 3, T, 8, 0, 3, T, 3, 0, 2, T, 6, 3, 4, 7

T marks the end of each time interval

During 1st interval

P	U3	U2	U1	U0
+	+	+	+	+
3	1	0	0	0
+	+	+	+	+
2	1	0	0	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+

At the end of 1st interval

P	U3	U2	U1	U0
+	+	+	+	+
3	0	1	0	0
+	+	+	+	+
2	0	1	0	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+

Additional Reference bits Algorithm

During 2nd interval

P	U3	U2	U1	U0
+---+	+---+	+---+	+---+	+---+
3	1	1	0	0
+---+	+---+	+---+	+---+	+---+
2	0	1	0	0
+---+	+---+	+---+	+---+	+---+
8	1	0	0	0
+---+	+---+	+---+	+---+	+---+
0	1	0	0	0
+---+	+---+	+---+	+---+	+---+
-	0	0	0	0
+---+	+---+	+---+	+---+	+---+

At the end of 2nd interval

P	U3	U2	U1	U0
+---+	+---+	+---+	+---+	+---+
3	0	1	1	0
+---+	+---+	+---+	+---+	+---+
2	0	0	1	0
+---+	+---+	+---+	+---+	+---+
8	0	1	0	0
+---+	+---+	+---+	+---+	+---+
0	0	1	0	0
+---+	+---+	+---+	+---+	+---+
-	0	0	0	0
+---+	+---+	+---+	+---+	+---+

Additional Reference bits Algorithm

During 3rd interval

P	U3	U2	U1	U0
+	+	+	+	+
3	1	1	1	0
+	+	+	+	+
2	1	0	1	0
+	+	+	+	+
8	0	1	0	0
+	+	+	+	+
0	1	1	0	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+

At the end of 3rd interval

P	U3	U2	U1	U0
+	+	+	+	+
3	0	1	1	1
+	+	+	+	+
2	0	1	0	1
+	+	+	+	+
8	0	0	1	0
+	+	+	+	+
0	0	1	1	0
+	+	+	+	+
-	0	0	0	0
+	+	+	+	+

Additional Reference bits Algorithm

During 4th interval, after 6, 3

P	U3	U2	U1	U0
+	+	+	+	+
3	1	1	1	1
+	+	+	+	+
2	0	1	0	1
+	+	+	+	+
8	0	0	1	0
+	+	+	+	+
0	0	1	1	0
+	+	+	+	+
6	1	0	0	0
+	+	+	+	+

During 4th interval, after 6, 3, 4

P	U3	U2	U1	U0
+	+	+	+	+
3	1	1	1	1
+	+	+	+	+
2	0	1	0	1
+	+	+	+	+
4	1	0	0	0
+	+	+	+	+
0	0	1	1	0
+	+	+	+	+
6	1	0	0	0
+	+	+	+	+

Page 4 replaced 8 with U bits 0010

Additional Reference bits Algorithm

During 4th interval, after 6, 3, 4, 7

P	U3	U2	U1	U0
3	1	1	1	1
7	1	0	0	0
4	1	0	0	0
0	0	1	1	0
6	1	0	0	0

If there are more victim pages with same U bits then chose FIFO

Page 7 replaced 2 with U bits 0101

Second Chance Algorithm

► Uses

- a reference/used bit
- NO history bits
- FIFO

Search FIFO pages circularly

If FIFO page's used bit is 0 then
 replace it

Else

 give that FIFO page a second chance
 clear its used bit
 goto next FIFO page

Second Chance Algorithm

P	U	3	P	U	2	P	U	3
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0 *			3 1			3 1		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0			0 *			2 1		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0			0			0 *		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0			0			0		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0			0			0		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		

Second Chance Algorithm

P	U	0	P	U	8	P	U	4
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
3 1			3 1			3 1		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
2 1			2 1			2 1		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0 *			0 1			0 1		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0			0 *			8 1		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - +		
0			0			0 *		
+ - - - + - - - +			+ - - - + - - - +			+ - - - + - - - -		

Second Chance Algorithm

P	U	2
+ - - - + - - - +		
3 1 *		
+ - - - + - - - +		
2 1		
+ - - - + - - - +		
0 1		
+ - - - + - - - +		
8 1		
+ - - - + - - - +		
4 1		
+ - - - + - - - +		

P	U	5
+ - - - + - - - +		
3 1 *		
+ - - - + - - - +		
2 1		
+ - - - + - - - +		
0 1		
+ - - - + - - - +		
8 1		
+ - - - + - - - +		
4 1		
+ - - - + - - - +		

P	U	0
+ - - - + - - - +		
5 1		
+ - - - + - - - +		
2 0 *		
+ - - - + - - - +		
0 0		
+ - - - + - - - +		
8 0		
+ - - - + - - - +		
4 0		
+ - - - + - - - +		

Second Chance Algorithm

P	U	9	P	U	8	P	U	3
+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +
5 1	5 1	5 1	5 1	5 1	5 1	5 1	5 1	5 1
+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +
2 0 *	2 0 *	2 0 *	9 1	9 1	9 1	9 1	9 1	9 1
+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +
0 1	0 1	0 1	0 1 *	0 1 *	0 1 *	0 1 *	0 1 *	0 1 *
+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +
8 0	8 0	8 0	8 0	8 0	8 0	8 1	8 1	8 1
+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +
4 0	4 0	4 0	4 0	4 0	4 0	4 0	4 0	4 0
+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +	+ - - - + - - - +

Second Chance Algorithm

P	U	2	P	U
+ - - - + - - - +	+ - - - + - - - +		+ - - - + - - - +	
5 1 *	5 0		5 0	
+ - - - + - - - +	+ - - - + - - - +		+ - - - + - - - +	
9 1	9 0		9 0	
+ - - - + - - - +	+ - - - + - - - +		+ - - - + - - - +	
0 0	2 1		2 1	
+ - - - + - - - +	+ - - - + - - - +		+ - - - + - - - +	
8 0	8 0 *		8 0 *	
+ - - - + - - - +	+ - - - + - - - +		+ - - - + - - - +	
3 1	3 1		3 1	
+ - - - + - - - +	+ - - - + - - - +		+ - - - + - - - +	

Enhanced Second Chance Algorithm

- ▶ Considers reference/used bit and dirty/modified bit
- ▶ Four cases, based on (R,M) , the pair:
 - ▶ $(R,M) = (0,0)$ neither recently used nor modified -- best to replace
 - ▶ $(R,M) = (0,1)$ not recently used but modified -- second choice
 - ▶ $(R,M) = (1,0)$ recently used but not modified -- third choice
 - ▶ $(R,M) = (1,1)$ recently used and modified -- fourth choice

Enhanced Second Chance Algorithm

- ▶ If class-1 page then use it
- ▶ If class-2 page then record first instance
- ▶ Else clear R bit and go to next FIFO page
- ▶ If scan is completed then
 - ▶ If class-2 page was recorded then use that recorded instance
 - ▶ Else repeat another scan

Counting based page replacements

Keep a counter of the number of references that have been made to each page

- ▶ **LFU Algorithm:** replaces page with smallest count
- ▶ **MFU Algorithm:** based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Allocation of Frames

- ▶ Each process needs *minimum* number of pages
 - ▶ Otherwise high page fault rate => slow execution
- ▶ Two major allocation schemes
 - ▶ Fixed allocation
 - ▶ Priority allocation

Fixed Allocation: Equal allocation

- ▶ Equal allocation – For example, if there are $m=100$ frames and $n=5$ processes, give each process 20 frames.
 - ▶ Smaller processes waste frames.
- ▶ Proportional allocation – Allocate according to the size of process

s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

Priority Allocation

- Use a proportional allocation scheme using priorities rather than size

Global vs. Local Allocation

- ▶ **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
- ▶ **Local replacement** – each process selects from only its own set of allocated frames
- ▶ In local replacement, no. of allocated frames is not affected by other processes => page fault rate is not affected by other processes.


Thrashing

- ▶ P_{high} steals frames from P_{low}
- ▶ Page fault rate increases for P_{low}
- ▶ P_{high} steals more frames from P_{low}
- ▶ Page fault rate further increases for P_{low}

- ▶ If a process does not have “enough” frames, the page-fault rate is very high. This leads to:
 - ▶ low CPU utilization
 - ▶ operating system thinks that it needs to increase the degree of multiprogramming
 - ▶ another process added to the system

- ▶ **Thrashing** \equiv a process is busy swapping pages in and out

Thrashing



Demand Paging and Thrashing

Why does demand paging work?

Locality model

- ▶ Process migrates from one locality to another
 - ▶ Currently executing **main** function (locality-1), may call **sqrt** function (locality-2)
- ▶ Localities may overlap
 - ▶ both the functions may use the same global pages

Why does thrashing occur?

- ▶ $\sum \text{size of locality} > \text{total memory size}$

Solution to Thrashing


- ▶ Use local page replacement algorithm.
- ▶ Use working-set model.
- ▶ Page fault frequency scheme.

Working-set Model


- ▶ $\Delta \equiv$ working-set window \equiv a fixed number of page references (say 10)
- ▶ **Working-Set** is the set of distinct pages in the working-set window.
- ▶ $WSS_i \equiv$ working set size of Process P_i

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$

Working-set Model

- ▶ Performance depends on Δ
 - ▶ if Δ too small will not encompass entire locality
 - ▶ if Δ too large will encompass several localities
 - ▶ if $\Delta = \infty \Rightarrow$ will encompass entire program
- ▶ $D = \sum WSS_i \equiv$ total demand frames
- ▶ if $D > m \Rightarrow$ Thrashing
- ▶ Policy if $D > m$, then suspend one of the processes

Working-Set model prevents thrashing while keeping the degree of multiprogramming as high as possible.

Page-Fault Frequency Scheme

- ▶ Establish “acceptable” page-fault rate
 - ▶ If actual rate too low, process loses frame
 - ▶ If actual rate too high, process gains frame

